


Whois N1: six models across feedback budgets

39 identical questions per point · one agent · one repetition · complete endpoint means



All $\beta=10$ points reuse historical results; dashed segments cross historical cohorts.

Gemini: $\beta=1/5/15/20$ via OpenRouter; $\beta=10$ via Sub2. No $\beta=10$ experiment was rerun.

Curves describe these execution cohorts; they do not isolate a budget-only causal effect.