


Whois N1: six models across feedback budgets

Protocol-repaired offline rescoring

39 identical questions per point · one agent · one repetition · all 1,170 scores known



$\beta=10$ reuses 234 main-experiment trajectories; dashed segments cross historical cohorts.

Gemini: $\beta=1/5/15/20$ via OpenRouter; $\beta=10$ via Sub2. All trajectories are existing executions.

Scores were repaired offline. These cohorts do not isolate a budget-only causal effect.